

MLBA Assignment 1 - Group 80

README

● Input Files:

The program takes paths to the training data as the user input. We have used the files provided on Kaggle as our training and testing data. For training data: train_data.csv For testing data: test_data.csv

● Output Files:

The program outputs a CSV file Group_80_Predictions.csv containing the predicted binary class labels (0,1) for the training data. The name of the prediction file submitted on Kaggle was different, though.

Prediction file submitted on Kaggle: cnn_att_3.csv

● Steps To Run the Python File:

Make sure that all dependencies and libraries used by this program are installed on your system.

A list of the required libraries are:

- Numpy
- Pandas
- Sklearn
- TensorFlow

Use the following steps to run the code:

Dependencies:

1. Sklearn
2. Tensorflow
3. Numpy
4. Pandas

Command Line Options:

1. Type the command "**python3 group80_code.py <train_csv_path> <test_csv_path>**" in the same directory as the notebook
2. Change the file paths to your own file paths. I have used the Google Drive paths as I saved my data there
3. Install all dependencies mentioned above.
4. After this, you can run all the cells u will get a CSV file as the output
"group80_predictions_output.csv"

Why Run?

To predict the labels for the protein sequences.

Model Info:

We used a CNN with attention. The CNN is a 1D CNN so it even checks the context of the neighboring Amino Acids by adding weights to it while making the embedding. Further, we use Self-attention to weigh the importance of each of the embedding going in. All these weights are learned by the model on its own. Finally, we add a few Dense Layers and an Output. We use Binary Cross entropy loss as the task is that of Binary Classification and the final layer has “sigmoid” activation.

Model Summary

Layer (type)	Output Shape	Param #	Connected to
input_13 (InputLayer)	[(None, None)]	0	[]
embedding_13 (Embedding)	(None, None, 16)	320	['input_13[0][0]', 'input_13[0][0]']
conv1d_39 (Conv1D)	(None, None, 100)	3300	['embedding_13[0][0]', 'embedding_13[1][0]']
attention_9 (Attention)	(None, None, 100)	0	['conv1d_39[0][0]', 'conv1d_39[1][0]']
global_average_pooling1d_10 (GlobalAveragePooling1D)	(None, 100)	0	['conv1d_39[0][0]']
global_average_pooling1d_11 (GlobalAveragePooling1D)	(None, 100)	0	['attention_9[0][0]']
concatenate_5 (Concatenate)	(None, 200)	0	['global_average_pooling1d_10[0][0]', 'global_average_pooling1d_11[0][0]']
dense_59 (Dense)	(None, 128)	25728	['concatenate_5[0][0]']
dense_60 (Dense)	(None, 64)	8256	['dense_59[0][0]']
dense_61 (Dense)	(None, 32)	2080	['dense_60[0][0]']
dense_62 (Dense)	(None, 16)	528	['dense_61[0][0]']
dense_63 (Dense)	(None, 1)	17	['dense_62[0][0]']

Run Summary

```
48/48 [=====] - 12s 245ms/step - loss: 0.3752 - accuracy: 0.8454
Epoch 32/40
48/48 [=====] - 12s 245ms/step - loss: 0.3697 - accuracy: 0.8514
Epoch 33/40
48/48 [=====] - 12s 245ms/step - loss: 0.3870 - accuracy: 0.8362
Epoch 34/40
48/48 [=====] - 12s 245ms/step - loss: 0.3793 - accuracy: 0.8448
Epoch 35/40
48/48 [=====] - 12s 245ms/step - loss: 0.3701 - accuracy: 0.8481
Epoch 36/40
48/48 [=====] - 12s 245ms/step - loss: 0.3702 - accuracy: 0.8487
Epoch 37/40
48/48 [=====] - 12s 246ms/step - loss: 0.3707 - accuracy: 0.8481
Epoch 38/40
48/48 [=====] - 12s 246ms/step - loss: 0.3759 - accuracy: 0.8461
Epoch 39/40
48/48 [=====] - 12s 245ms/step - loss: 0.3709 - accuracy: 0.8408
Epoch 40/40
48/48 [=====] - 12s 245ms/step - loss: 0.3791 - accuracy: 0.8428
<keras.src.callbacks.History at 0x7ba716e25900>
```

Further models tried:

1. Convolutional Neural Network + Attention
2. Random Forest
3. Transformer
4. Easy Classifier
5. Bagging

The best accuracy given was **Convolutional Neural Network + Attention**.